

Claude Token Usage: Executive Brief

DataLoaderPlatform | 2026-08-24 | For internal engineering alignment

STATIC PROMPT OVERHEAD

~24.5k

API DOCS CORPUS

~76.5k

DESIGN DOCS CORPUS

~144.4k

SAVINGS OPPORTUNITY

45%-80%

What is driving cost

- Default routing to multi-stage agent chains.
- Repeated reloading of long global instructions.
- Multiple agents re-reading large API/design docs.
- Review and test loops with broad context each cycle.

What to change

- Adopt fast lane by default for non-schema changes.
- Reserve full lane for new loaders and shape changes.
- Compress prompts to checklist format.
- Run scoped tests first; escalate only when needed.

Expected token savings


New loader builds: 45% to 65%



Existing loader changes: 60% to 80%



Review/test loops: 30% to 50%



Execution readiness

The implementation playbook is already prepared and ready to run immediately after loader delivery.

- docs/superpowers/plans/2026-08-24-claude-token-usage-optimization-playbook.md

Prepared by GitHub Copilot (GPT-5.3-Codex). Method: repository instruction audit and measured-size token estimation.